

DDL

Grigorios Loukides
Email: grigorios.loukides@kcl.ac.uk

Session Objectives

Data
Definition
Language
2/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

In this session, you will learn:

- How to translate ER models into tables in a database
- Structured Query Language (SQL)
- Data types supported by SQL standard
- How to create, alter and delete tables
- How to create and delete indexes using SQL

Database Design Problem

Data
Definition
Language
3/42

Draw an ERD for next week:

The club has a number of members who can practice different martial arts. For all members you need to register, in addition to the name and the membership number, which martial arts they practice (you can practice several) and which belt (or degree) they hold in the art in question. You also need to register information about their membership fee, namely the amount and the payment date. You must even register whether the member has a valid licence or not (you need a different licence for each art that you practice, so think about how to represent these!)

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

From ERD to Relational Model

Data
Definition
Language
4/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

ERD	Logical Model
Entity	Create table
Attribute	Create column in table
Relationship	
-1:N	Primary key of table on "one" side as foreign key of table on "many" side
-1:1	
(a) Mandatory on both sides	Combine entities into one table
(b) Mandatory on one side	Primary key of table on "optional" side as foreign key of table on "mandatory" side
(c) Optional on both sides	Arbitrary
-N:M relationship	Create a composite table. The primary keys from entities into the new table to act as primary and foreign key

Introduction

Data
Definition
Language
5/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

SQL is a database language that:

- Allows you to create database and table structures, to perform data management tasks and to perform complex queries designed to transform the raw data into useful information.
- It is portable, it is a de facto standard SQL

SQL functions fit into two broad categories:

- Data Definition Language (DDL): SQL includes commands to create database objects such as tables, indexes, and views, as well as commands to define access rights to those database objects
- It is a data manipulation language (DML): SQL includes commands to insert, update, delete, and retrieve data within the database tables

Data Definition Language (DDL)

DDL allows to define database objects:

- SCHEMA is a collection of tables
- TABLE is a set of data elements (values) using a model of vertical columns (identifiable by name) and horizontal rows, the cell being the unit where a row and column intersect. A table has a specified number of columns, but can have any number of rows
- INDEX is a special lookup table that the DBMS can use to speed up data retrieval
- VIEW is a virtual table that does not necessarily exist in the database but can be produced upon request

Writing SQL

Data
Definition
Language
8/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

SQL statement consists of *reserved words* and *user-defined words*:

- *Reserved words* are a fixed part of SQL and must be spelt exactly as required
- *User-defined words* are made up by user and represent names of various database objects such as relations, columns, views

Literals in SQL

Data
Definition
Language
9/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

- Literals are constants that are used in SQL statements
- There are different forms of literals for every data type supported by SQL

Writing Literals

A literal is a data value commonly used in variable assignments or comparisons:

- *Numeric Literals*: Integers are represented as a sequence of digits. Floats use . as a decimal separator
- *String Literals*: is a sequence of bytes or characters, enclosed within either single quote (') or double quote (") characters
- *Date and Time Literals*: Date and time values can be represented in several formats, such as quoted strings or as numbers, depending on the exact type of the value and other factors
- *Boolean Literals*: The constants TRUE and FALSE evaluate to 1 and 0, respectively. The constant names can be written in any lettercase.

Syntax Notation

Data
Definition
Language
11/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

- Upper-case letters represent reserved words
- | indicates a choice among alternatives.
- Curly braces { indicate a required element.
- Square brackets [indicate an optional element.
- ... indicates optional repetition (0 or more).

Overview of some DDL Commands

Data
Definition
Language
12/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

CREATE SCHEMA	DROP SCHEMA
CREATE TABLE	DROP TABLE
ALTER TABLE	
CREATE INDEX	DROP INDEX
CREATE VIEW	DROP VIEW

Creating a Database

```
CREATE {DATABASE | SCHEMA} [IF NOT EXISTS] dbName
```

Creates a database with the given name:

- To use this statement, you need the CREATE privilege for the database.
- CREATE SCHEMA is a synonym for CREATE DATABASE.
- An error occurs if the database exists and you did not specify IF NOT EXISTS.

Deleting a Database

```
DROP {DATABASE | SCHEMA} [IF EXISTS] dbName
```

Deletes a database with the given name:

- Be very careful with this statement
- To use this statement, you need the DROP privilege for the database.

Create Table

Data
Definition
Language
15/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

```
CREATE TABLE tableName
{
  (colName dataType [NOT NULL] [DEFAULT defaultOption]
    [AUTO_INCREMENT] [,...])
  [PRIMARY KEY (listOfColumns),]
  {[FOREIGN KEY (listOfFKColumns)
    REFERENCES parentTableName [(listOfCKColumns)] [,])
  ]}
```

Creates a table with the given name

Data Types

Data
Definition
Language
16/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

MySQL uses many different data types, broken into three categories:

- 1 Numeric
- 2 Date and time
- 3 String types

Numeric Data Types

- INT - A normal-sized integer that can be signed or unsigned
 - TINYINT, SMALLINT, MEDIUMINT, BIGINT
- FLOAT(M,D) - A floating-point number that cannot be unsigned. You can define the display length (M) and the number of decimals (D). This is not required and will default to 10,2, where 2 is the number of decimals and 10 is the total number of digits (including decimals)
 - DOUBLE (M,D), DECIMAL(M,D)

Date and Time Types

- DATE - A date in YYYY-MM-DD format
- TIME - Stores the time in HH:MM:SS format
- DATETIME - A date and time combination in YYYY-MM-DD HH:MM:SS format

String Types

- **CHAR(M)** - A fixed-length string between 1 and 255 characters in length, right-padded with spaces to the specified length when stored. Defining a length is not required, but the default is 1
- **VARCHAR(M)** - A variable-length string between 1 and 255 characters in length. You must define a length when creating a VARCHAR field
- **ENUM** - An enumeration, which is a fancy term for list. When defining an ENUM, you are creating a list of items from which the value must be selected (or it can be NULL). For example, if you wanted your field to contain "A" or "B" or "C", you would define your ENUM as ENUM ('A', 'B', 'C') and only those values (or NULL) could ever populate that field.

Exercise

Data
Definition
Language
20/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

Given the table below determine the data types of the different columns:

Table name: EMPLOYEE

EMP_NUM	EMP_LNAME	EMP_FNAME	EMP_INITIAL	EMP_HIREDATE	JOB_CODE
101	News	John	G	08-Nov-00	502
102	Senior	David	H	12-Jul-89	501
103	Arbough	June	E	01-Dec-96	503
104	Ramoras	Anne	K	15-Nov-87	501
105	Johnson	Alice	K	01-Feb-93	502

Create Table

Data
Definition
Language
21/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

```
CREATE TABLE tableName
{ (colName dataType [NOT NULL] [DEFAULT defaultOption]
  [AUTO_INCREMENT] [, ...] }
[PRIMARY KEY (listOfColumns),]
{ [FOREIGN KEY (listOfFKColumns)
  REFERENCES parentTableName [(listOfCKColumns)] [, ,] }
  ) }
```

- The optional DEFAULT clause can be specified to provide a default value for a particular column
- The AUTO_INCREMENT attribute can be used to generate a unique identity for new rows
 - If no value is specified for an AUTO_INCREMENT column, MySQL assigns sequence numbers automatically
 - If applies to integer and floating-point types

Create Table

Data
Definition
Language
22/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

```
CREATE TABLE tableName
{ (colName dataType [NOT NULL] [DEFAULT defaultOption]
  [AUTO_INCREMENT] [, ...] }
[PRIMARY KEY (listOfColumns),]
{ [FOREIGN KEY (listOfFKColumns)
  REFERENCES parentTableName [(listOfCKColumns)] [,] }
  ) }
```

- The PRIMARY KEY clause specifies the column or columns that form the primary key
- The FOREIGN KEY clause specifies a foreign key

Create Table

Data
Definition
Language
23/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

```
CREATE TABLE tableName
{ (colName dataType [NOT NULL] [DEFAULT defaultOption]
[,...] }
[PRIMARY KEY (listOfColumns),]
{ [FOREIGN KEY (listOfFKColumns)
REFERENCES parentTableName (listOfCKColumns) [,] }
)} }
```

- The FOREIGN KEY clause specifies a foreign key
 - A listOfFKColumns, the column or columns from the table being created that form the foreign key.
 - A REFERENCES subclause, giving the parent table; that is, the table holding the matching candidate key and the list of columns referenced

Exercise

Data
Definition
Language
24/42

Logical
Database
Design

SQL

DDL

Database
Management
**Table
Management**
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

Given the attributes below (which are a subset of the Employee table) create a SQL statement to create the table of a table named EMP_1

ATTRIBUTE (FIELD) NAME	DATA DECLARATION
EMP_NUM	CHAR(3)
EMP_LNAME	VARCHAR(15)
EMP_FNAME	VARCHAR(15)
EMP_INITIAL	CHAR(1)
EMP_HIREDATE	DATE
JOB_CODE	CHAR(3)

Foreign Key

Data
Definition
Language
25/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

```
FOREIGN KEY (listOfFKColumns)
REFERENCES parentTable (listOfFKColumns)
    [ON DELETE option]
    [ON UPDATE option]
```

option: RESTRICT | CASCADE | SET NULL | SET DEFAULT

MySQL:

- Rejects any INSERT or UPDATE on the child table violating referential integrity
- An UPDATE or DELETE operation in a value in the parent table with matching rows in the child table is determined by the referential action

Foreign Key

```
FOREIGN KEY (listOfFKColumns)
REFERENCES parentTable (listOfFKColumns)
    [ON DELETE option]
    [ON UPDATE option]
option: RESTRICT | CASCADE | SET NULL | SET DEFAULT
```

Referential Actions:

- **CASCADE:** Delete or update the row from the parent table, and automatically delete or update the matching rows in the child table
- **SET NULL:** Delete or update the row from the parent table, and set the foreign key column or columns in the child table to NULL
- **RESTRICT:** Rejects the operation for the parent table
- For an ON DELETE or ON UPDATE that is not specified, the default action is always RESTRICT

Cloning a Table

Data
Definition
Language
27/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

```
CREATE TABLE newTable LIKE oldTable
```

Creates an empty table based on the definition of another table, including any column attributes and indexes defined in the original table

Altering a Table

Data
Definition
Language
28/42

ALTER TABLE *tblName* alterSpecification

Changes the structure of a table

- These changes are indicated in the **alterSpecification**
- You can add or delete columns, keys, indexes, change the type of existing columns, or rename columns or the table itself..
- More information about its syntax:
<http://dev.mysql.com/doc/refman/5.7/en/alter-table.html>

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

Deleting a Table

```
DROP TABLE [IF EXISTS] tblName[,...]
```

DROP TABLE removes one or more tables (whose names are passed as a list):

- All table data and the table definition are removed, so be careful with this statement!
- If any of the tables named in the argument list do not exist, MySQL returns an error indicating by name which nonexistent tables it was unable to drop, but it also drops all of the tables in the list that do exist.

Exercise

Data
Definition
Language
30/42

Table name: DIRECTOR Database name: Ch03_Theater

DIR_NUM	DIR_LNAME	DIR_DOB
100	Broadway	12-Jan-65
101	Hollywoody	18-Nov-53
102	Goofy	21-Jun-62

Table name: PLAY

PLAY_CODE	PLAY_NAME	DIR_NUM
1001	Cat On a Cold, Bare Roof	102
1002	Hold the Mayo, Pass the Bread	101
1003	I Never Promised You Coffee	102
1004	Silly Putty Goes To Washington	100
1005	See No Sound, Hear No Sight	101
1006	Starstruck in Bloxi	102
1007	Stranger In Parrot Ice	101

- 1 What are the data types of the different columns?
- 2 What are the referential integrity constraints that should hold on the schema?
- 3 Write appropriate SQL DDL statements to define the database

Index

Data
Definition
Language
31/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

- Indexes are used to find rows with specific column values quickly.
- Without an index, a DBMS must begin with the first row and then read through the entire table to find the relevant rows.
- The larger the table, the more this costs.
- If the table has an index for the columns in question, DBMS can quickly determine the position to seek to in the middle of the data file without having to look at all the data.
- This is much faster than reading every row sequentially.

Index Example

Data
Definition
Language
32/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

staffNo	fName	lName	salary	branchNo		salary
SG14	David	Ford	18000	B003	↗	9000
SG37	Ann	Beech	12000	B003	←	12000
SL21	John	White	30000	B005	↖	18000
SL41	Julie	Lee	9000	B005	↘	30000

Index

Data
Definition
Language
33/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

- An index is a structure that provides accelerated access to the rows of a table based on the values of one or more columns
- Indexes can be used to improve the efficiency of data retrievals
- However, since indexes may be updated by the system every time the underlying tables are updated, additional overheads may be incurred
- Indexes are usually created to satisfy particular search criteria after the table has been in use for some time and has grown in size.

Create Index

Data
Definition
Language
34/42

```
CREATE [UNIQUE] INDEX indexName ON tableName  
    (columnName [ASC| DESC][, . . . ])
```

- The specified columns constitute the index key and should be listed in major to minor order
- If the UNIQUE clause is used, uniqueness of the indexed column or combination of columns will be enforced by the DBMS
- For each column, we may specify that the order is ascending (ASC) or descending (DESC)

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion
Suggested
Readings

Index Design

Data
Definition
Language
35/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

- It's not possible to make meaningful physical design decisions until you understand in detail the operations that have to be supported by the DBMS
- General Recommendations
 - 1 Do not index small tables
 - 2 Index the primary key (MySQL does it automatically)
 - 3 Add an index to any column that is heavily used for data retrieval
 - 4 Add a secondary index to a foreign key if there is frequent access based on it

Drop Index

Data
Definition
Language
36/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

```
DROP INDEX indexName ON tableName
```

Drops the index named *indexName* from the table *tableName*

Exercise 1

Data
Definition
Language
37/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

Hotel (hotelNo, hotelName, city)
Room (roomNo, hotelNo, type, price)
Booking (hotelNo, guestNo, dateFrom, dateTo, roomNo)
Guest (guestNo, guestName, guestAddress)

- 1 Create the Hotel table
- 2 Now create the Room, Booking, and Guest tables with the following constraints:
 - type must be one of Single, Double, or Family.

Exercise 2

Data
Definition
Language
38/42

Logical
Database
Design

SQL

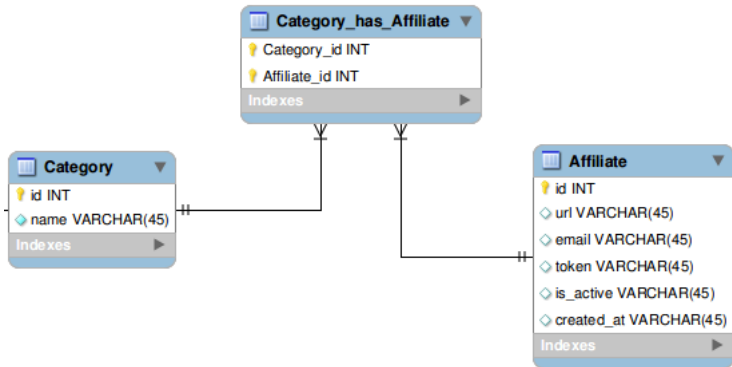
DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings



- 1 What are the referential integrity constraints that should hold on the schema?
- 2 Write appropriate SQL DDL statements to define the database

Exercise 3

*A company needs to store information about employees (identified by *ssn*, with *salary* and *phone*), departments (identified by *dno*, with *dname* and *budget*), and children of employees (with *name* and *age*). Employees work in departments; each department is managed by an employee; a child must be identified uniquely by name when the parent (who is an employee; assume that only one parent works for the company) is known. We are not interested in information about a child once the parent leaves the company.*

- 1 Drawn an ERD for these requirements
- 2 Write appropriate SQL DDL statements to define the database

Conclusion

Data
Definition
Language
40/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

In this session we have covered:

- Table creation, alteration and deletion
- Index creation and deletion
- **Feedback on KEATS** - this week for lecturers, next week for labs/tutorials.

Lab Session

**Data
Definition
Language**
41/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

Next week's lab session more about DDL in MySQL

Suggested Readings

Data
Definition
Language
42/42

Logical
Database
Design

SQL

DDL

Database
Management
Table
Management
Index
Management

Tutorial
Exercises

Conclusion

Suggested
Readings

- Chapters 5 of Fundamentals of Database Systems. Elmasri & Navathe.
- Chapter 6 of Database systems: a practical approach to design, implementation, and management. Connolly, Thomas M; Begg, Carolyn